Inference often combines conditioning and summarizing

Observed		**Condition**		**Summarize**
y	$\rightarrow$	$p(z \mid y) = \mathcal{N}\left(z \mid \mu_{z\mid y}, \Sigma_{z\mid y}\right)$	$\rightarrow$	$\hat{z} = E[z \mid y] = \mu_{z\mid y}$

Murphy, Book 2 §2.3.2.2–2.3.2.3

$$E[z \mid y] = \int z p(z \mid y)\, \mathrm{d}z$$

- Conditioning preserves the whole posterior distribution and its uncertainty
- Summarizing extracts a quantity from that posterior; for a Gaussian the mean is $\mu_{z\mid y}$

Choose a reconstruction by its conditional squared-error risk

Observed y; choose $a \in \mathbb{R}^d$ to reconstruct unknown $z \in \mathbb{R}^d$

Squared-error criterion

$$L(a, z) = \|z - a\|^2, \quad R(a|y) = E[L(a, z)|y]$$

Center at the conditional mean

$$m = E[z|y], \quad r = z - m, \quad z - a = r + (m - a)$$

$$\|z - a\|^2 = \|r\|^2 + 2r^T(m - a) + \|m - a\|^2$$

- Choose a to minimize conditional expected squared error $R(a|y)$.
- Given y, offset $m - a$ is fixed; uncertainty remains in r.

Centering the error makes the posterior mean optimal

Average; $m - a$ fixed Previous slide: pointwise squared-error expansion

$$\|z - a\|^2 = \|r\|^2 + 2r^T(m - a) + \|m - a\|^2$$

$$E[r|y] = E[z|y] - m = 0 \qquad V = \mathrm{Cov}[z|y]$$

$$E[r^T(m - a)|y] = E[r|y]^T(m - a) = 0 \qquad E[\|r\|^2|y] = \sum_{j=1}^{d} E[r_j^2|y] = \mathrm{tr}\, V$$

Conditional squared-error risk

$$R(a|y) = \mathrm{tr}\, V + \|m - a\|^2$$

$$\|m - a\|^2 \geq 0 \quad \text{equality iff} \quad a = m$$

$$a^*(y) = m = E[z|y]$$

$$\min_a R(a|y) = \mathrm{tr}\, V$$

$$\text{Earlier Gaussian:} \quad a^*(y) = \mu_{z|y}$$

$$\min_a R(a|y) = \mathrm{tr}\, \Sigma_{z|y}$$

- Risk is residual uncertainty plus squared distance from the posterior mean.
- The posterior mean minimizes model expected squared error; covariance measures what remains.

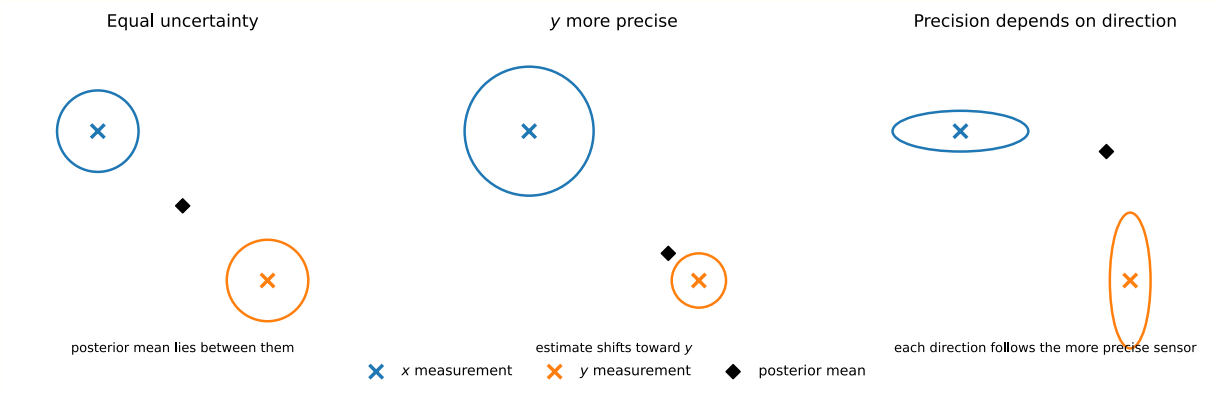
Sensor fusion combines measurements according to their uncertainty

Two noisy measurements of one state

$$z \sim \mathcal{N}(\mu_z, \Sigma_z)\; x \mid z \sim \mathcal{N}(z, \Sigma_x), \quad y \mid z \sim \mathcal{N}(z, \Sigma_y)$$

Same two operations

$$p(z \mid x, y) \quad \rightarrow \quad E[z \mid x, y]$$



Course schematic after Murphy, Book 2 §2.3.2.3, Fig. 2.10; `sensor_fusion_figure.ipynb`

- Equal uncertainty puts the posterior mean between the measurements
- A more precise sensor pulls the posterior mean closer to its measurement
- With anisotropic uncertainty, the weighting can differ by direction